

Database - Collection of related data

A database represents some aspect of real world

DBMS - Database Management System

Software to create, maintain and access data

Physical Level (Internal):

Physical Data model, data storage, access path

Logical Level (conceptual):

Describes entire database: entities, data types, relationships, constraints
Physical & data independence

View Level (External):

describes the part of the database that a particular user group is interested in and hides rest.

~~Data isolation - Data separated from queries~~

Database Architecture

- 1) Centralized - Located and stored in a single location
- 2) client server - Users connect to the database through network

- 3) Parallel - Multiple processors work in parallel on the database to provide resources.
 - 4) Distributed - Collection of multiple interconnected databases, which are spread physically across various locations that communicate via computer network.

Domain types in SQL

- Superkey is an attribute that will identify any tuple must be unique

Relational Operators Algebra

- Selection - selecting rows - $\sigma_{B=17}^{(1)}$ → match rows where B column is 17
- Projection - selecting columns - $\Pi_{A,C}^{(1)}$ → remove duplicate rows
- Joining - Two or more tables columns
↳ If the table cartesian product - it forms all possible combinations $\rightarrow T_1 \text{ rows} \times T_2 \text{ rows} = \text{Total rows}$
- Union - combines 2 same columns from 2 different tables and only retains unique rows.
- Difference (R-S) → Tuple should be in R but not in S
- Intersection → Tuples common in 2 tables (1)
- Rename (P) - Rename a column while performing cartesian product
- Natural Join - It joins tuple where the common column value matches
 - i) Left outer join - keep all from left and matching from right
 - ii) Right outer join - keep from right
 - iii) Outer join - Only ~~keep~~ common columns.

SQL

entity integrity - not null, unique
foreign key references other tables

ex: create table instructor (

ID char(5),
name varchar(20) not null,
dept_name varchar(20),
salary numeric(8,2),
primary key (ID),
foreign key(dept_name) references dept_name)

* reference table;
reference key should be created first

* primary key declaration ensures attribute not null.

ex: drop table instructor - deletes instructor table only rows not structure

delete table instructor - Only deletes rows table structure remains intact.

alter table instructor (add grade varchar(2))
to add column.
to drop use drop in alter table construction.

SQL Query - select A₁, A₂, ..., A_n
from n₁, n₂, ..., n_m
where P.

String Operations

Like operator is used for pattern checking

- i) % - It matches any substring
- ii) underscore (-) - It matches character

ex: _intro% looks for string beginning with intro
--- looks for 3 char string

'100\%' used to compare substring including %.

Concatenation ("")

Pattern matching is case sensitive

Ex: Select SUBSTR('A B C D E F', 2, 3) → B C D
↳ 3 chars from 2

Select * from instructor, teacher
↳ cartesian product

Select ID, name, salary/12 as monthly salary

Compare multiple columns

where (instructor.ID, def.name) = (teacher.id, "C")

· Select course_id from selection where sem='Fall' and year=2009.

AND → Union operation

Null values

· The result of any arithmetic operation on null is null.

$$5 + \text{NULL} = \text{NULL}$$

Degree of Relationship set

- Binary relationships - involves two entity sets

- Relationships between more than 2 entity sets
are rare - ~~most~~ They are called ternary

Attributes

- An entity is represented by a set of attributes

- Domain - The set of permitted values of each attribute

Attribute types:

Having more than one
members

i) Single & Composite

ii) Single valued and multi valued (Phone nos)

iii) Derived attributes - Can be computed
from other attributes

e.g. Age from DOB

Mapping Cardinality Constraints

Expresses the number of entities to which
another entity can be associated via a relationship
set

- Types of mapping cardinality for binary relationships

if

i) one to one

ii) one to many

iii) Many to one

iv) Many to many

- CHEN's notation to be used
 - Rectangles - Entity sets
 - Diamonds - Relationship sets
 - Underline - Primary key
 - Attributes listed inside entity rectangle - One
 - One-cardinality
 - Many-cardinality
- Discriminator or Partial key underlined with a dashed line
- Entity is a key which has different attributes associated with it.
- ★ If object can be expressed by its properties then it's an entity.
- ★ A relationship is an association among several entities.
ex: 4453 (Austin) advisor 22222 (Einstein)
Student

Attribute Types:

- Simple and composite attribute
- Single-valued and multivalued attributes
- Derived attributes

Symbols

E → Entity set

R → Relationship set

R → Identifying relationship set for weak entity set

R E → Total participation of entity set in
Total participation of relationship
Total participation of values in E

E
A ₁
A ₂
A _{2.1}
A _{2.2}
of A ₃
A ₄ ()

Attributes:

Simple (A₁)

(A₁)

Composite (A₂) and

(A₂)

Multivalued (A₂)

Derived (A₄)

(A₄)

E
A ₁

Primary Key

E
...A ₁ ...

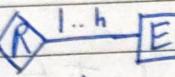
Discriminating attribute of
weak entity set

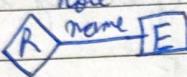
→ They don't have primary key
and can't be identified on their own.

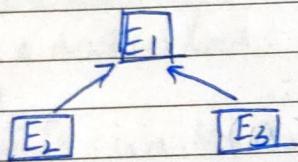
 many-to-many relationship

 many-to-one relationship

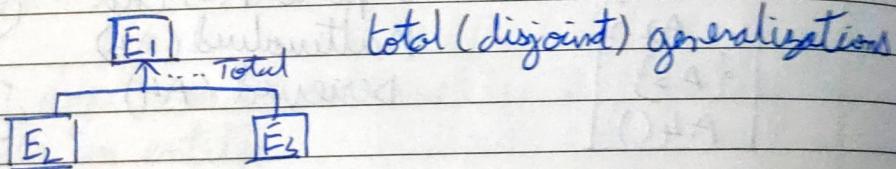
 one-to-one relationship

 $R \xrightarrow{1..n} E$ Cardinality limits

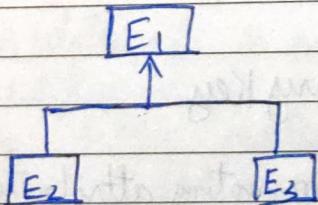
 $R \xrightarrow{\text{role-name}} E$ Role indicator



ISA: Generalization or specialization



disjoint generalization



Intermediate SQL

NULL and Three Value

Any comparison with NULL = NULL

Three Value Logic

OR : Unknown or True = True

" or False = Unknown

" or Unknown = Unknown

AND : Unknown and True = Unknown

" and False = False

" and Unknown = Unknown

NOT : (Not Unknown) = Unknown

Aggregate Functions

avg, min, max, sum, count

Using group by :

```
select dept_name, avg(salary) as avg_salary  
from instructor  
group by dept_name;
```

gives aggregate for each dept name

HAVING clause with group by.

To get dept_name for having avg_salary > 4200

select dept_name, avg(salary) as avg_salary
from instructor
group by dept_name
having avg(salary) > 42000

- * Having is used for statements having group by
- * Where is used for general queries

Nested Subqueries

A subquery is a select-from-where expression that is nested within another query.

Set Membership:

Find courses offered in Fall 2017 and Spring 2018

select distinct course_id
from section
where semester = 'Fall' and year = 2017 and
course_id in (select course_id
from section
where semester = 'Spring' and year = 2018);

Not can be added to check opposite case

'Some' clause

iii. Find names of instructors who with salary greater than that of some (at least one) instructor in Bio dept.

select name
from instructor
where salary > some (select salary
from instructor
where dept_name = 'Biology')

'All' clause

Replace some with all salary value greater than others

Test for Empty Relations

The exists construct returns the value true if the argument subquery is non empty

exists \Rightarrow true if $r \neq \emptyset$
not exists \Rightarrow true if $r = \emptyset$

Except can be used to add exceptions to exists clause.

Unique used to check if duplicate tuple is present in table.

'with' clause (*Not supported in many DBMS)

It provides a way of defining a temporary relation whose definition is available only to the query in which the with clause occurs.

Scalar Subquery

List A subquery which returns only one value.

Runtime error if more than one tuple returned

Deletion

delete all from instructor : deletes all instruction

· Delete all instructor where from Finance Department
Delete from instructor where dept_name = 'V';

Insertion

Ex: Make each student in Music department who earned more than 144 credit hours an instructor in the Music department with salary of \$180.00

Insert into Instructor

Select ID, name, dept_name, 18000

from Student

where dept_name = 'Music' and total_credits > 144;

Update

qn.i) Give a 5% salary hike to all instructor

Update instructor

$$\text{set salary} = \text{salary} * 1.05$$

ii) Give a 5% raise to those lesser than average

Update instructor

$$\text{set salary} = \text{salary} * 1.05$$

when ~~(select salary avg(salary))~~

$\text{salary} < (\text{select avg salary from instructor})$,

Case Statement

Ex.i) Hike 5% for less than 100000 and hike 3% ~~if~~ $\text{key} > 100000$

Update instructor

$$\text{set salary} = \text{case}$$

when $\text{salary} < 100000$ then $\text{salary} * 1.05$

else $\text{salary} * 1.03$

end.

* Not equal \Leftrightarrow

Natural or Conditional Joins

Natural join - joins by matching column

Natural left outer join - join by matching but keeps all tuples from left table

1/1/8 for right outer join

Inner join → course inner join ~~as~~ prereq on
course . course_id = prereq . course_id.
→ Joined on condition column is repeated

Outer join retains left table tuples

Views

Temporary variables not stored such as tables created using Select

i) View of instructors without salary

create view faculty as select ID, name, dept_name
from instructor

ii) To create custom columns

create view departments_total_salary (dept_name,
total_salary) as select dept_name, sum(salary)
from instructor group by dept_name;

view defined using other views

create view physics-fall-2009 as select course-course-id, sec-section-id, building, room-number from course, section where course.course-id=section.course-id and course.dept_name='Physics' and section.semester='Fall' and section.year='2009'

To find courses only of Watson

create view physics-fall-2009-watson as select course-id, room-number from physics-fall-2009 where building='Watson'

Insert statement is same when inserted in view
It inserts in parent table also missing value takes up NULL value.

* A view is created with values of history department
insertion will fail if another subject is tried to be inserted

To prevent the insertion check option should be used.

Materialized View - Table created with all the tuples present in parent table

Integrity constraints on single relation

Not null

Unique:

- i) Unique specification allows the values to form a candidate key.
- ii) Candidate keys can be null (whereas primary key can't)

Check (P)

Referential Integrity

To show when value is updated in one table referenced table should be updated

Concrete table "X"

foreign key (dept_name) references department
on delete cascade
on update cascade,

)

Complex check clauses

create assertion <name> check (& predicates)

Built-in data types

Date '2005-07-27'

Time: '09:00:00.75'

Timestamp: '2005-7-27 09:00:00.75'

Interval: $\text{en: Interval} \text{ '1' day}$

User-defined types

e.g. create type Rupee as numeric (12, 2)

Domain

Data type with constraints

ex: create domain ~~constraint~~ person-name char(20)
not null.

Large-Object types

blob - binary large object

clob - character large object

Authorization Specification in SQL

- The grant statement is used to confer authorization.

grant <privilege list> on
<relation name or view name> to <user list>

User list:

- a user id
- public
- A role

e.g.: Grant select on instructor to U₁, U₂, U₃
" all on " to give all

Revoking

- The revoke statement is used to revoke authorization.
revoke <privilege list>
on <relation name or view name> from <user list>

e.g. revoke selection
revoke select on branch from U₁, U₂, U₃

Role

create role instructor;

grant instructor to Amit;

↳ To add give someone the role

Roles given to other roles

create role teaching-assistant

grant teaching-assistant to instructor

Transfer of privileges

grant select on department to Amit with grant option

revoke select on department from Amit, Satoshi cascade

↳ If Amit permission is removed whoever he granted will also lose

revoke select on department from Amit, Satoshi restrict

↳ Opposite of cascade

grant reference (dept_name) on department to Mariano

If dept_name is present in another table which he is editing he can also edit table in referenced table.

PL SQL

- Procedural Language extensions to SQL
- SQL is a popular language for both querying and updating data
- PL/SQL adds many procedural constructs to SQL language to overcome some limitations of SQL.

PLSQL Block

Declaration Section

Begin

Execution Section

Exception

Execution Section

END;

i) DECLARE

I-Customer-group VARCHAR2(100); := 'Silver';
BEGIN

I-Customer_group := 'Gold';

DBMS_OUTPUT.PUT_LINE(I-(Customer-group));

END;

ii) DECLARE

JNAME Instructor_Name%Type;

JNAME gets same data type as name in Instructor

Begin

Select more from instructor when id = '1, 2, 3';
DBMS_OUTPUT.PUT_LINE('Name : ' || Salary);

END;

If then statement

DECLARE

m_sales NUMBER := 2000000;

BEGIN

IF m_sales > 100000 THEN

DBMS_OUTPUT.PUT_LINE('Sales revenue
is greater than 100k');

END IF;

END;

Case Statement

... Declare ...

Begin

c_grade := 'B';

CASE c_grade

WHEN 'A' THEN

c_rank := 'Excellent'

WHEN 'B' THEN

c_rank := 'Very Good';

ELSE

c_rank := 'No such grade';

END CASE;

DBMS_OUTPUT.PUT_LINE(c_rank);

END

LOOP

```
BEGIN
    LOOP
        I_Counter := I_Counter + 1
        IF I_Counter > 3 THEN
            EXIT;
        END IF;
        DBMS_OUTPUT.PUT_LINE('... || I_Counter');
    END LOOP;
    -- Control resumes here after exit
END
```

FORLOOP

```
BEGIN
    FOR I_Counter IN 1..5 OR WHILE m_Counter <= 8
    LOOP
        DBMS_OUTPUT.PUT_LINE(I_Counter);
        m_Counter := m_Counter + 1
    END LOOP;
END;
```

EXIT when counter = 3
↳ acts as break

CONTINUE done as usual

Procedure (function)

It is a reusable unit that encapsulates specific business logic of that application.

It is a named block stored as schema object in Oracle database.

Ex: CREATE [OR REPLACE] PROCEDURE procedure-name
(parameters-list) IS [Declaration statements]
BEGIN
[Execution statements]
EXCEPTION
[Exception Handlers]
END [Procedure Name];

Parameters

- i) IN - Read only, default, can be referenced but value can't be changed
- ii) OUT - Writable parameters, used for output
- iii) INOUT - Readable and writable parameters

If not specified considered as IN

Function

- Almost same as procedure return statement is needed

```
CREATE OR REPLACE FUNCTION deft_count(  
    deft_name IN varchar)  
RETURN number IS  
    de_count number := 0  
     $\hookrightarrow$  declaration of data type
```

- Function must always return something
- Functions can be used in typical SQL statements like select, Update, Insert etc.

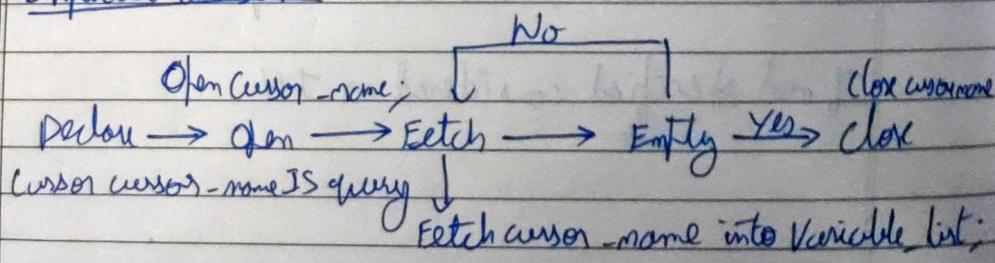
Cursor

A pointer that points to a result of a query.

Implicit cursor:

When select into, insert, update and delete addition automatically creates implicit cursors.

Explicit cursor:



→ Refers to instructor table
cursor cinstri IS select * From Instructor

De-instr = c_instri%ROWTYPE,

begin

OPEN c_instri;

LOOP

FETCH c_instri INTO n_instr;

EXIT WHEN c_instri%NOT FOUND;

If n_instr.salary < 60000 THEN

Update instructor set salary = salary*10
where id = instr.id

END IF;

DBMS_Output.Put_line ('*' .. *)

ENDLOOP

CLOSE c_instri

END;

To be put inside a procedure or function

// Start of ER

- Binary relationship - involve two entity sets
- An entity is represented by a set of attributes
- Domains - set of permitted values for each attribute

Normalization

- Process to get schema into proper form
- Normal forms 1NF, 2NF, 3NF, BCNF

Functional Dependencies

α and β and β'

$$\alpha \rightarrow \beta$$

For any tuples t_1 & t_2

$$t_1[\alpha] = t_2[\alpha] \text{ then } t_1[\beta] = t_2[\beta]$$

e.g.:

	α	β	
t_1	1	4	
t_2	1	5	$\therefore \alpha \rightarrow \beta$ does not hold
t_3	3	7	True

Key - Value of other attribute can be uniquely identified using a key.

Superkey - Identify all attributes in relation

Candidate Key - Minimal can't be broken down
 $K \rightarrow R$ for which K is a candidate key then
 $\alpha \rightarrow R$ and $\alpha \subset K$ then $\alpha \rightarrow R$
 α should not uniquely identify R

$\alpha \rightarrow \beta$ is trivial ; $\beta \subseteq \alpha$

(I , D , $None$) $\rightarrow ID$
 $\alpha \rightarrow \beta$
 $\beta \subseteq \alpha$

	EmpID	Name	P-ID	D-ID
t1	1	JD	1	HR
t2	2	Ine	2	Marketing
t3	3	John	1	HN
t4	4	Global Glocal	3	Sales

Check if $F: \{ EmpID \rightarrow Name \}$ holds true
 $\alpha \rightarrow \beta$

As they are all unique it holds true

$$F: \left\{ \begin{array}{l} D-ID \rightarrow D-None \\ \alpha \rightarrow \beta \end{array} \right\} \quad \begin{array}{l} t_1[1] = HR = t_1[1] \\ t_1[1] = HR \\ t_1[1] = HR \end{array}$$

: Holds true

* Closure of a set of functional dependencies (F^+)

$A \rightarrow B$ and $B \rightarrow C$ then $A \rightarrow BC$
 i.e., f.d. infers f.d. $\rightarrow F^+$

Armstrong's Rules

Additional

R1 if $\alpha \subseteq BC$ then $\alpha \rightarrow P$ (reflexivity) R4: $\alpha \rightarrow \beta$ and $\alpha \rightarrow \gamma$
 then $\alpha \rightarrow \beta \gamma$ (union)

R2 if $\alpha \rightarrow \beta$ then $\delta\alpha \rightarrow \delta\beta$ (nugation)

R5: $\alpha \rightarrow \beta$ then $\alpha \rightarrow \beta$

R3 if $\alpha \rightarrow \beta$ and $\beta \rightarrow \gamma$ we have $\alpha \rightarrow \gamma$
 (transitivity)

and $\alpha \rightarrow \gamma$

(decomposition)

R6: $\alpha \rightarrow \beta$ then $\gamma\beta \rightarrow \gamma$

then $\gamma\alpha \rightarrow \gamma$

(Pseudotransitivity)

Ex: $R = (A, B, C, D, E, F, G, H, I)$

$F: \{ AB \rightarrow E, AG \rightarrow I, BE \rightarrow I, E \rightarrow G, GJ \rightarrow H \}$

Can $AB \rightarrow GH$ be inferred?

$AB \rightarrow E, E \rightarrow G \Rightarrow AB \rightarrow GH$

$AB \rightarrow BE = AB \rightarrow BE$

~~AB~~

$AB \rightarrow BE \rightarrow J \Rightarrow AB \rightarrow I. — (6)$

$AB \rightarrow E$ and $E \rightarrow G \Rightarrow AB \rightarrow G - (7)$

F from (6) and (7)

$AB \rightarrow IG$ (union) - (8)

we have $GI \rightarrow H \Rightarrow AB \rightarrow H - (9)$

from (7) and (9)

By union $AB \rightarrow GH$

$$F^+ = \{ F, 6, 7, 8, 9, 10 \}$$

$$2) F = \{ A \rightarrow BC \\ (D \rightarrow E) \\ E \rightarrow C \\ D \rightarrow AEH \\ ABH \rightarrow BD \\ DH \rightarrow BC \}$$

$$\text{or } R = \{ A, B, C, D, E, H \}$$

Show ~~BCD~~ $BCD \rightarrow H$ hold

$$(D \not\rightarrow E, D \rightarrow AEH, CD \rightarrow CAEH)$$

$$BCD \Rightarrow BCAEH$$

By decomposition

$BCD \rightarrow H$ (decomposition)

3) $A \rightarrow BC$
 $B \rightarrow E$
 $CD \rightarrow EF$
ST $AD \rightarrow F$

$CD \rightarrow EF \Rightarrow D \rightarrow$

~~$AD \rightarrow BC$~~ $A \rightarrow BC \Rightarrow A \rightarrow B$ and $A \rightarrow C$ (Q)
 ~~$B \rightarrow D$~~ ~~$B \rightarrow F$~~
 ~~$AD \rightarrow BEF$~~ $AD \rightarrow CD - \textcircled{S}$

From \textcircled{S} $(D \rightarrow EF \Rightarrow CD \rightarrow F - \textcircled{G})$

From \textcircled{S} $A \textcircled{H}$

Get $AD \rightarrow F$

$F = \{$
 $A \rightarrow BC$
 $B \rightarrow E$
 $CD \rightarrow EF$
 $A \rightarrow R$
 $A \rightarrow C$
 $AD \rightarrow CD$
 ~~$A \rightarrow E$~~ $(D \rightarrow F)$
 $CD \rightarrow F$
 $AD \rightarrow F$

Attribute Closures

α under F

α^+ set of all attributes that can determine F

If $\alpha \rightarrow \beta$; $\beta \subseteq \alpha$

Ex: $R(A, B, C, G, H, I)$

$A \rightarrow B$, $A \rightarrow C$, $G \rightarrow H$, $GA \rightarrow I$, $B \rightarrow H$

Find $(AG)^+ = \{AG, \underbrace{B, C, H, I}_{\text{as } AG \rightarrow AG} \}$

$\hookrightarrow \text{using } G \rightarrow H$

$\cap A \rightarrow B, A \rightarrow C$

∴ Final $(AG)^+ = AGBCHI$

Uses

- 1) Check if FDs hold
- 2) Find Candidate key and Superkey
PK
- 3) Prime attribute and non prime attribute

Ex1: $R(A, B, C, D, E, F, G)$

fds : $AB \rightarrow CD$, $AF \rightarrow P$, $DE \rightarrow F$, $C \rightarrow G$, $F \rightarrow E$,
 $G \rightarrow A$.

Do $CF \rightarrow DF$ hold?

$CF^+ \subseteq CEG \sqsupseteq A \sqsupseteq D$

$\therefore CF \rightarrow DF$ is true

Ex 2) Find the candidate keys and super keys

$R(A, B, C, D, E)$

fds: $AB \rightarrow C$, $DE \rightarrow B$, $CD \rightarrow E$

~~* By seeing FDS as AD can be seen as key on left side
so start with them~~

$$A^+ = A$$

$$D^+ = D$$

$$AD^+ = AD$$

$$ADB^+ = ADB \xrightarrow{B} AB \rightarrow C$$

$$\hookrightarrow A \cancel{C} \cancel{E} \mid A \cancel{D} \cancel{B} \cancel{C} \cancel{E}$$

$$ADC^+ = ADC \sqsupseteq B$$

$$ADE^+ = ADE \sqsupseteq B \sqsupseteq C$$

Note if $\alpha \rightarrow R$; α is a candidate key

In this case ADB^+ , ADC^+ , ADE^+ as they
define all. Any one of these candidate keys can
be primary key.

Superkeys :- Superset of candidate keys

$ADB, ADBC, ADBE, AD BCE$
 $ADC, ADCE, \dots$

Similarly all can be formed.

Prime attributes: PA = {A, B, C, D, E}

Non prime - ?

AO all are prime

Q2: R(A, B, C, D, E, F)
 $AB \rightarrow C, CF \rightarrow D, D \rightarrow E, F \rightarrow B, E \rightarrow F$

Find candidate, superkeys, prime attributes
 no prime

$A^+ = A \rightarrow$ If this is candidate key don't add more letters

~~AB~~ = ABCDEF

~~ABC~~ = ABCDEF

~~ABD~~ = ABCDEF

~~ABE~~ = ABECFD

~~AC~~ = ACDEFBD

~~AD~~ = ADEFBC

~~AF~~ = AEFBCD

~~AF~~ = AFBCDE

Candidate keys: AB, AC, AD, AF, AE

Superkeys = AB, ABC, ABCD, ABCE, ABCF, ABCD, ABCF

Prime Attributes = PA = {A, B, C, D, E, F}

Non prime = { }
G

Canonical Cover

Ex1: $A \rightarrow B$ $B \rightarrow C$ $A \rightarrow C$ ← redundant fd

Ex2: $A \rightarrow B, B \rightarrow C, A \rightarrow CD$
 $\hookrightarrow A \rightarrow G$ $\circlearrowright A \rightarrow D \rightarrow$ more essential
 Redundant

Entourage Attribute

$\alpha \rightarrow \beta$ in F

i) α as entourage in α and F logically implies
 $F = (\alpha \rightarrow \beta) \cup (\alpha - A) \rightarrow \beta$

ii) ~~$\alpha \rightarrow \beta$~~

ii) A as entourage in β

$(F - (\alpha \rightarrow \beta)) \cup \alpha \rightarrow (\beta - A)$ logically imply F

'A Test for extraneous attribute

1) If $A \in \alpha$

i) complete $(\alpha - A)^+$ using F

~~LHS~~ ii) check if $(\alpha - A)^+$ contains β , then A is extraneous

2) If $A \in \beta$

i) complete α^+ using $F' = F - (\alpha \rightarrow \beta) \cup \alpha \rightarrow (\beta - A)$

~~RHS~~ iii) check α^+ contains A, if so A is extraneous

Ex1) R(A, B, C)

$A \rightarrow BC$ - ①

$B \rightarrow C$ - ②

$AB \rightarrow C$ - ③ $A \rightarrow B$ - ④

Find the canonical cover

i) Redundant factors

ii) Extraneous attribute removal

iii) Step 1 is repeated

A) From ① & ④ we observe ④ can be inferred from ①, so 4 is redundant

Now set

$A \rightarrow BC$ - ①

$B \rightarrow C$ - ②

$AB \rightarrow C$ - ③

Checking (1) & (2) for extraneous

$$\begin{array}{c} A \rightarrow B \\ A B \rightarrow C \end{array}$$

check if A is extraneous

check if B is extraneous

$$\text{Check } (A - A)^T = B$$

$$B^+ = BC$$

Even after removing A, BC is obtained
∴ A is extraneous.

Check if B is extraneous - no check needed as
at least one attribute needed.

Now (2) becomes $B \rightarrow C$ is available as
② hence should be removed.

$$\text{Now } f_2 = A \rightarrow BC \quad - (1)$$

$$B \rightarrow C \quad - (2)$$

Check extraneous in attribute in (1) for PHS

$$A^+ = AC \text{ for } P \Rightarrow A \rightarrow C$$

\hookrightarrow As B is not present here it is not
extraneous

$$P_2 \text{ for } C: A \rightarrow B$$

$$B \rightarrow C$$

$A^+ = ABC$ and C is available so C is
extraneous.

Find canonical cover

$$A \rightarrow B$$

$$B \rightarrow C$$

Ques 2
R(A, B, C, D)

$$fd = \{ A B \rightarrow C D, BC \rightarrow D \}$$

Find canonical cover

A.) Step 1: check for redundant fds:
No redundant fd.

Step 2: Check extraneous
if present remove

Equivalent or not

Find G are fds

F can be inferred using G and
 G can be inferred using F

then $F = G$

G covers F : If all fds of G can be inferred by F .

G covers F : If all fds of F can be inferred by G .

F covers G and G covers F then
 F equivalent to G .

$$\text{Q1: } F: \begin{cases} A \rightarrow C & -\textcircled{1} \\ AC \rightarrow D & -\textcircled{2} \\ E \rightarrow AD & -\textcircled{3} \\ F \rightarrow H & -\textcircled{4} \end{cases} \quad G: \begin{cases} A \rightarrow CD \\ E \rightarrow AH \end{cases}$$

If G covers F using fds in G

$A^+ = ACD \rightarrow \textcircled{1} \& \textcircled{2}$ of F is covered

$E^+ = EAH \rightarrow \textcircled{4}$ of F is inferred

G covers F

If F comes in

Using (1) of F
At AC

Using (2) of F

Act - ACD \rightarrow (1) of g is inferred

Using (3) of F

E + EAD \rightarrow Using (4) of F

EAD H \rightarrow (2) of a is inferred

Lossless Join Decomposition

To identify lossless

- $R_1 \cup R_2 = R$
- $R_1 \cap R_2 \neq \emptyset$
- $R_1 \cap R_2 \rightarrow R_1 \text{ or } R_1 \cap R_2 \rightarrow R_2$

Ex: $R = \{A, B, C\}$

$F = \{A \rightarrow B, B \rightarrow C\}$

can be decomposed in 2 different ways

$$R_1 = (A, B), R_2 = (B, C)$$

Lossless join

$$R_1 \cap R_2 = \{B\} \text{ and } B^+ - BC$$

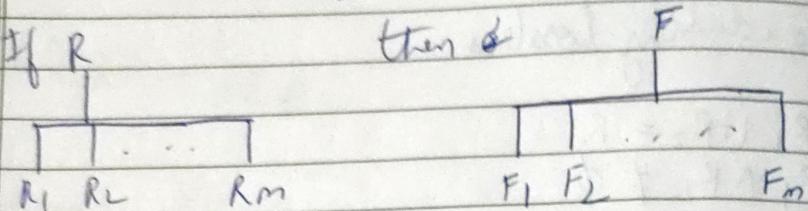
$$R_1 = (A, B) \cdot R_2 = (A, C)$$

$$R_1 \cap R_2 = \{A\} \text{ and } A^+ - AB$$

Not dependency preserving

(cannot check $B \rightarrow C$ without computing $R_1 \Delta R_2$)

Dependency Preserving



A decomposition is dependency preserving if
 $(F_1 \cup F_2 \cup F_3 \dots \cup F_m) = F^+$

Common attribute should be used as subkey
in order to preserve dependencies

To check dependency preserving:

$$\text{1) } F_1 \cup F_2 = F$$

$$\text{e.g. } R(ABCD, EF)$$

$$F: A \rightarrow BCD, A \rightarrow EF, BC \rightarrow AD, BC \rightarrow E, BC \rightarrow F, B \rightarrow F \\ D \rightarrow EF$$

→ New tables.

$$D = \{ABC, D, BF, DE\}$$

F_1	F_2	F_3
<u>ABC(D/R1)</u>	<u>BF(R2)</u>	<u>DE(R3)</u>
$A \rightarrow BCD$	$B \rightarrow F$	$D \rightarrow E$
$BC \rightarrow AD$		

Need to check for $A \rightarrow BCD$, $(A \rightarrow EF)$, $BC \rightarrow AD$

$$(BC \rightarrow E), (BC \rightarrow F), B \rightarrow F, D \rightarrow E$$

only one not in table need to be found

$$(BC)^+ F_1 = (ABCD)^+ F_2 = (B, RADF) | F_3 = BCADFE$$

$$(A^+)F_1 \cdot ABCD)^T F_2 = \{ABCDF\} F_3 = ABCDFE$$

: It is dependency preserving

Normal Forms

MPurpose :

- Eliminating redundant (useless) data
- Ensuring data dependencies make sense.

Anomalies

1) Update Anomaly - Deconform the schema

Employee table having ID, address and skill.

If employee has multiple address and one skill.

multiple skills then everytime skill is address is updated it might not reflect everywhere.

better to deconform into (IP,Address), (ID,skill)

- 2) Insertion Anomaly, value can't be entered into table due to null value or other constraint
- 3) Deletion constraint if faculty has no course assigned his value would get deleted.

1NF

- All relations only contain Atomic values
- Atomic values - Single valued attributes
- This form has lot of redundancies

2NF

• Checks:

→ It should be 1NF

→ Should not contain Partial Dependency

A value part of candidate key if it defines non prime attribute it is not in 2NF

3NF

- Should be in 2NF
- It should not contain transitive dependencies

$A \rightarrow B$ $B \rightarrow C \Rightarrow A \rightarrow C$ Should be avoided

Converting into 3NF

Given: relation R, set F of functional dependencies

Find: Decomposition of R into a set of 3NF relations
 R_i

Steps:

1) Eliminate ~~redundant~~ redundant FDs, resulting in a canonical cover F_c of F

2) Create a relation $R_i = X$ for each FD $X \rightarrow x$ in F_c

3) If the key K of R doesn't occur in any ~~other~~ relation R_i , create one more relation $N_i = K$

BCNF Boyce Codd Normal Form

- Doesn't guarantee ~~less~~ gain and dependency preserving
- Guarantees lossless join

R(N, D, P, S, A)

$N \rightarrow D$: fd₁
 $S \rightarrow A$: fd₂
 $(NS) \rightarrow P$: fd₃

a. What is highest possible normal form/3NF relation
 to 2NF, 3NF, BCNF

i) Check for 2NF

Candidate Key = (N, S)

fd₁ = Violation of 2NF as part of it identifies
 individually non prime attribute

$N \rightarrow D$

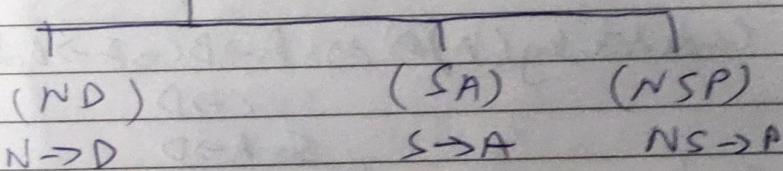
~~S → A~~

fd₂ = Not in 2NF

fd₃ = In 2NF

∴ R is not in 2NF

N, D, P, S, A



is in 3NF

As all dependency mapped so dependency preserving.

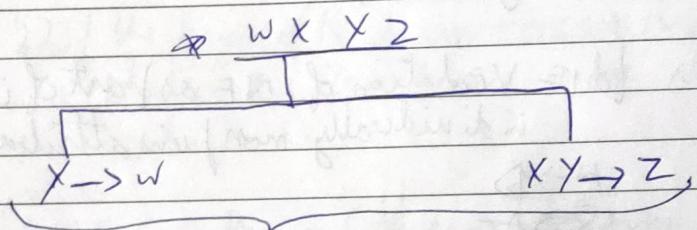
Qn2: $R(w, x, y, z)$
 $XY \rightarrow z - \text{fd}_1$
 $y \rightarrow w - \text{fd}_2$

Candidate key (X, Y , ~~w~~)

$\text{fd}_1 \rightarrow \cancel{\text{2NF}}$ $\text{Im } 2\text{NF}$

$\text{fd}_2 \rightarrow \text{not } 2\text{NF}$ as individually mapping to w

R not 2NF



As only 1 fd satisfies 3NF and dependency preserving

Qn3: $R(A B C D)$

$A \rightarrow BC$

$C \rightarrow D$

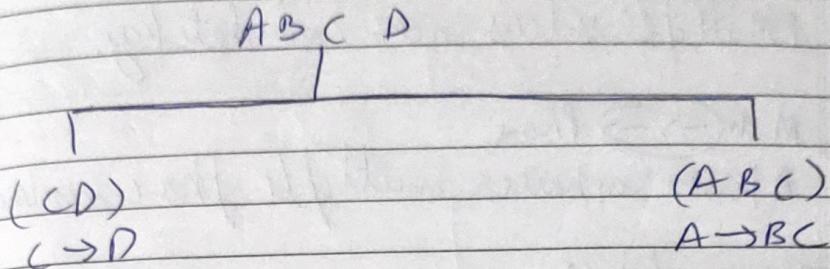
Candidate key (A) as $A \rightarrow BC \Rightarrow A \rightarrow B$ and $A \rightarrow C$
 $C \rightarrow D$

$\therefore A \rightarrow D$

Therefore as $C \rightarrow D$ & C is not candidate key

It is 2NF

f_{d2} violates $\leq N F$ as $NPA \rightarrow NPA$



Q. $R(CSJD P QV)$

$f_{d1} \Rightarrow C \rightarrow SS \oplus SDPQV$

$f_{d2} = SD \rightarrow P -$

$f_{d3} \Rightarrow JP \rightarrow C$

$f_{d4} \Rightarrow J \rightarrow S$

Candidate key is C , so f_{d2}, f_{d3}, f_{d4} are not in $3NF$

$A J \rightarrow S - (1)$

$SD \rightarrow R - (2)$

$JP \rightarrow P$ can be obtained

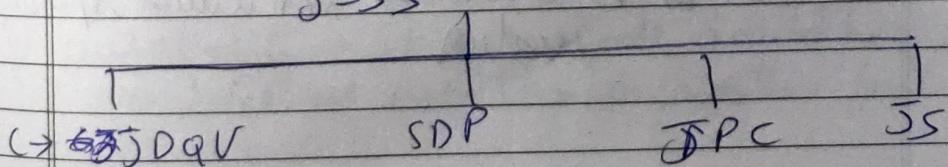
For each f_d in CC make a relation if key is not in my decomposed table then create new relation having key.

$CC \Rightarrow C \rightarrow J D Q V$

$SD \rightarrow P$

$JP \rightarrow C$

$J \rightarrow S$



Only one dependency in each case so $BCNF$

Multivalued Dependency

FM multiple values in one candidate key.

Man $\rightarrow \rightarrow$ Phone

A person can have multiple phone numbers.

Theory of MVDs

Complementation: $\neg (x \rightarrow \rightarrow y)$ then $x \rightarrow \rightarrow R - \{x, y\}$

Augmentation: $x \rightarrow \rightarrow y$ and $w \geq 2$ then $wx \rightarrow \rightarrow yz$

Transitivity: $x \rightarrow \rightarrow y$ and $y \rightarrow \rightarrow z$ then
 $x \rightarrow \rightarrow (z - y)$

Replication: $x \rightarrow \rightarrow y$ then $x \rightarrow \rightarrow y$ but reverse not true

Co-absence: $x \rightarrow \rightarrow y$ and there is w such that
 $w \cap y$ is empty, $w \geq 2$ and $y \geq 2$ then
 $x \rightarrow \rightarrow z$

A MVD $x \rightarrow \rightarrow y$ in R is trivial if

y is a subset of x ($x \geq y$) or

$x \cup y = R$ otherwise non-trivial repeating values

Transactions

A Series of one or more SQL statements that are logically related are termed as Transaction.

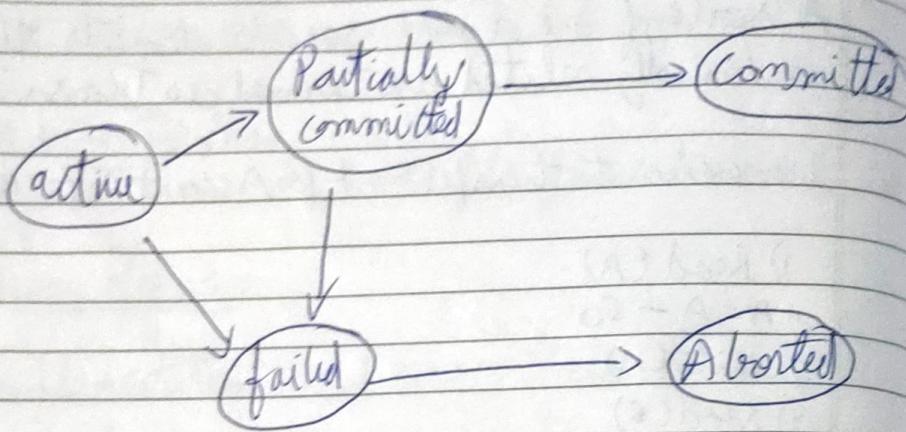
Ex: Transaction to transfer 50 \$ from Account A to B

- 1) Read(A)
- 2) $A = A - 50$
- 3) Write(A)
- 4) Read(B)
- 5) $B = B + 50$
- 6) Write(B)

* ACID properties

- 1) Atomicity - All operations should be completely executed or none
- 2) Consistency - Execution of a transaction in isolation preserves consistency of data base.
Ex: When 50 dollars transferred to B then B must be incremented
- 3) Isolation - Multiple executions taking place on a singular database. Transactions should be conducted in isolation
- 4) Durability - Changes made should persist in database even though system fails.

Transaction Model



Commit and Rollback commands

- A transaction begins with first executable SQL statement after a commit, rollback or connection made to a database.
- Transaction can be closed using commit or rollback.
- Commit ends current transaction and makes permanent changes
- Rollback ends transaction but undoes any changes made during transaction.

Syntax

commit;

rollback;

Savepoint marks and saves the current point in the processing of a transaction.

Savepoint deposit:

Rollback to Savepoint deposit;

The Lost Update Problem

This occurs when two transactions that access the same database have their operations interleaved in a way that makes the value of some database item incorrect.

The Temporary Update Problem (Dirty Read)

This occurs when one transaction updates a database item and then transaction fails for some reason.

Incorrect Summary Problem

If one transaction is calculating an aggregate summary function on a number of records while other transactions are updating some of these records.

Schedules

Schedule - A sequence of instructions that specify the chronological order in which the instructions of concurrent transactions are executed.

A schedule S of transactions T_1, T_2, \dots, T_m is an ordering of all operations in these transaction.

i.e. given

$$T_1 = R_1(Q) W_1(Q) \quad T_2 = R_2(Q) W_2(Q)$$

the order must be maintained

- a schedule $R_1(Q) R_2(Q) W_1(Q) W_2(Q)$

concurrent schedules used in real life situations.

A concurrent schedule is serializable if it is equivalent to a serial schedule.

Serializability methods:

- 1) conflict Serializable
- 2) View Serializable

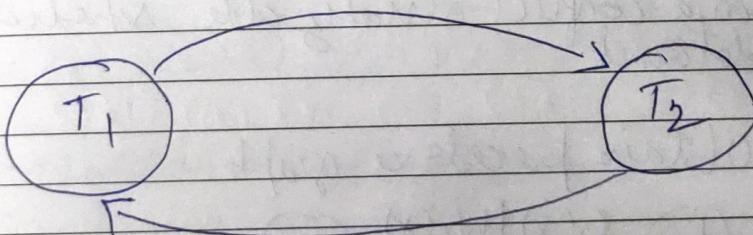
Conflict serializability

1) ~~Def of t_i = Real~~ (a)

If any schedule s can be transformed into a schedule s' by a series of swaps of non conflicting instructions, we say that s and s' are conflict equivalent.

We say schedule s is conflict serializable if it is conflict equivalent to a serial schedule.

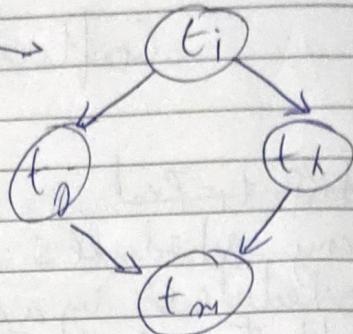
Precedence Graph - A directed graph where vertices are transactions



A schedule is conflict serializable if and only if its precedence graph is acyclic

If precedence graph is acyclic, the serializable order can be obtained by a topological sorting of the graph.

Topological sort \rightarrow



- 1) Take node with no incoming edges
and remove t_i
 $O/P \Rightarrow t_i$

- 2) Repeat same
 $O/P: t_i, t_j$

- 3) O/P: t_i, t_j, t_k, t_m

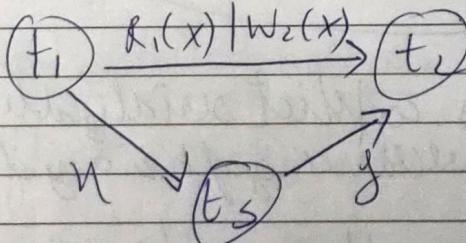
There can be more than one orders

(*)

$$\begin{array}{c} R_1(x) \quad R_2(y) \quad R_3(y) \quad W_2(y) \quad W_1(x) \\ \downarrow \quad \downarrow \quad \downarrow \\ W_3(x) \quad R_2(x) \quad N_2(x) \end{array}$$

Con a conflict serializable schedule be obtained.

Step 1: Obtain precedence graph



No loop so no conflict

4(b) 1 chronological seqn

$$t_1 \rightarrow t_2 \rightarrow t_3$$

Soln: $R_1(X) W_1(X) R_3(Y) W_3(X) R_2(Y)$
 $W_2(X) R_2(X) W_2(X)$

Q. $\overbrace{R_1(A)}^{R_2} W_1(A) \quad R_2(A) \overbrace{W_2(A)}^{R_3} \quad \overbrace{R_1(B)}^{R_2} W_1(B) \quad R_2(B) \overbrace{W_2(B)}^{R_3}$

$$(t_1) \longrightarrow (t_2) \quad t_1 \rightarrow t_2$$

Soln: $R_1(A) W_1(A) R_2(B) W_2(B) R_3(B) W_3(B)$
 $R_1(B) W_1(B)$

Ans: $\overbrace{W_1(X)}^{W_2(X)} \quad \overbrace{W_2(Y)}^{W_3(Y)} \quad W_3(Y)$

No precedence graph cyclic

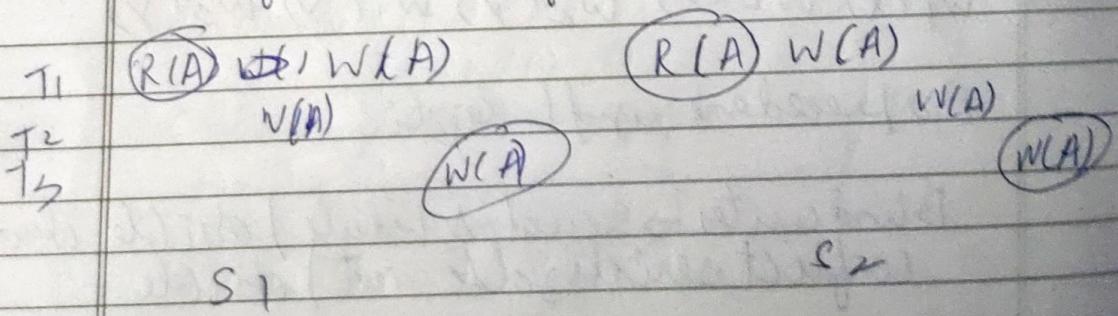
Blind writes \rightarrow serial schedule possible but conflict serializable not possible.

When conflicts, serializable not possible
try view serializable schedule

View Serializability

Schedules s_1 and s_2 are view equivalent if:

- T_i reads initial value of A in s_1 , then T_i also reads initial value of A in s_2
- If T_i reads value of A written by T_j in s_1 , then T_i also reads value of A written by T_j in s_2
- If T_i reads final value of A in s_1 , then T_i also writes final value of A in s_2

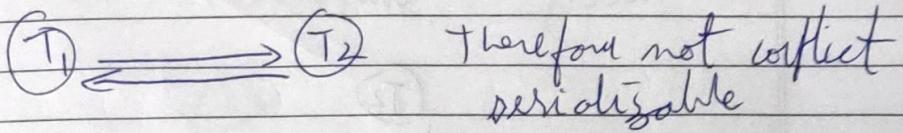


- Every conflict serializable is also view serializable
- Every view serializable that is not conflict serializable is Period writer.

Qn: $R_1(A) \quad W_1(A) \quad R_2(A) \quad W_2(A)$

If it is view serializable

If yes give its view serializable



$\xrightarrow{T_3}$ no update before A

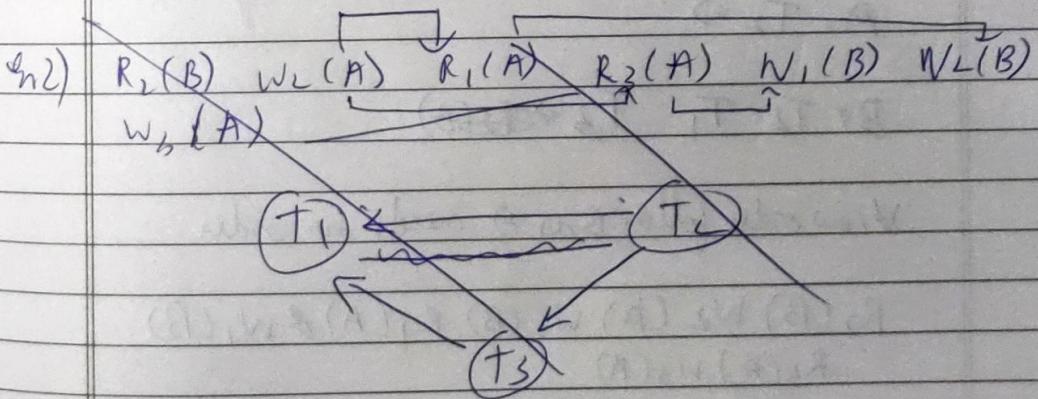
\rightarrow means no update before

Step:	Initial read	T_1
	Final Update	T_3
	All writes	$T_2 \quad T_1 \quad T_3$

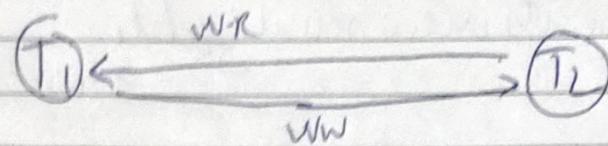
Read has to take place first

A: $T_1 \quad T_2 \quad T_3 \rightarrow$ Final write

In blue remaining operations
 $\Rightarrow R_1(A) \quad W_1(A) \quad R_2(A) \quad W_3(A)$



$R_2(B) W_2(A) \underbrace{R_1(A)}_{WR} R_3(A) \underbrace{W_1(B)}_{WW} \underbrace{W_2(B)}_{WR} W_3(B)$



$\overset{\text{not}}{(T_2)}$
Blind writes so conflict serializable

step 2:

	A	B
Initial Read	R X	R T ₂
Final Update	W T ₂	W T ₂
all writes	T₁ T ₃	T ₁ T ₂ T ₃
	Random ↓ only write	↑ same as your real should be first

$\cancel{W_1(B) R_2(B) W_2(A) \cancel{R} W_3(B) R_1(A) \cancel{W} T_2(B)}$
 $\cancel{R_3(A) W_3(A)}$

A; T₂ \Rightarrow

B; T₂ T₁ T₃ \Rightarrow ~~T₂(B)~~

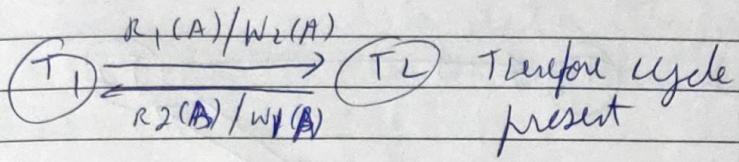
View order follows B as A needs no order:

$R_2(B) W_2(A) W_2(B) R_1(A) \cancel{R} W_1(B)$
 $R_3(A) W_3(A)$

3)

$$S : R_1(A) \ R_2(A) \ W_2(A) \ R_2(B) \ W_1(A) \ R_1(B) \ W_1(B)$$

Step 1: Procedure

~~(T2)~~

Step 2:

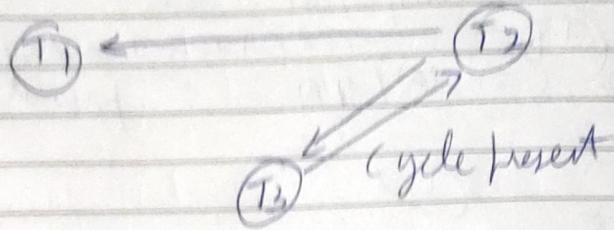
	A	B
Initial	T_1	T_2
Final Update	T_1	T_2 or T_1
All write	$T_2 - T_1$	T_1

No change before R_1 reads

A: $T_2 < T_1$
 B: $T_2 - T_1$

↓ serial order

4) $R_1(A) R_2(A) R_3(B) W_1(A) \swarrow R_2(C) R_2(B) W_2(B)$
 $W_3(C)$



	A	B	C
Initial Read	$T_1 T_2$	$T_3 T_2$	T_2
Final Update	T_1	T_2	T_3
All writes	T_1	T_2	T_3

A: ~~$T_2 \rightarrow T_1$~~ $T_2 \rightarrow X \rightarrow T_1$
 B: $T_3 \rightarrow Y \rightarrow T_2$ } ordering conflicting
 C: $T_2 \rightarrow X \rightarrow T_3$

So view serializable schedule possible.

Recoverable Schedules

- If a transaction T_j reads a data item previously written by transaction T_i , then the commit operation of T_i appears before commit of T_j .

Cascading Rollback

- A single transaction failure leads to a series of transaction rollbacks.

Corcodeless Schedules

- For each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation T_i appears before the read operation T_j .

Levels of Consistency in SQL 92

- Serializable - default
- Repeatable Read - Only committed records to be read, repeated reads must return same value
- Read Committed - Only committed records can be read but successive reads of record may return different.
- Read Uncommitted - Even uncommitted records may be read.

Concurrency Control

Lock Based Protocol

Lock mechanism

exclusive (X) for write operations
shared (+S) for read operations

lock - X instruction

lock - S instruction

Lock compatibility matrix

		T ₂	
		S	X
T ₁	S	True	False
	X	False	False

Any number of transactions can hold shared locks on an item

If any transaction holds an exclusive on the item no other transaction may hold any lock on the item

Two Phase Locking Protocol

Phase 1: Growing phase

→ Transaction may obtain locks

Phase 2: Shrinking phase

→ Transaction may release locks.

Protocol ensures serializability.

Cascading Roll Back solution

- i) Strict two-phase locking:

Here a transaction must hold all its ~~exclusive~~^{abs} locks till it commits / aborts. After commit it must unlock.

- ii) Rigorous two-phase locking is even stricter.
Here all locks are held ~~till~~ commit / abort
X or S

Dependency - Multiple transactions asking for locks when another transaction has done.

Deadlock Handling

Deadlock Prevention protocol ensures that system will never enter into a deadlock state

Strategies:

Fairness to abort

- i) Wait-die scheme (non preemptive)

Older transaction may wait for younger ones to release data item. Younger transaction never wait for older ones they are rolled back instead.

i) Round-visit-preemption

- Older transaction rounds off younger transaction instead of waiting for it.
Younger transaction may wait for older only.

(a) Suppose in a database T_1, T_2 and T_3 with timestamps 10, 20, 30 respectively. T_1 is holding a data item which T_1 and T_3 are requesting to acquire. T_1 come with respect to round visit

$T_1 \rightarrow$ continue

$T_2 \rightarrow$ Abort

$T_3 \rightarrow$ wait for T_2

wrt Wait-dil

$T_1 \rightarrow$ Wait for T_2

$T_3 \rightarrow$ Abort

Timestamp-Based protocol

The protocol maintains for each data 'Q' two time stamp values:

- $W\text{-tstamp}(Q)$ is the largest time-stamp of any transaction that executed write(Q) successfully.
 - $R\text{-tstamp}(Q)$ is the largest time-stamp of any transaction that executed read(Q) successfully.
- Timestamp ordering protocols ensure that any conflicting read and write operations are executed in time stamp order.

Eg.: For transactions 1, 2, 3, 4, 5. Apply timestamp based protocol.

	1	2	3	4	5	
T ₁	T ₁	T ₂	T ₃	T ₄	T ₅	
Read(Y)		Read(Y)				
		Write(Y)				
			Write(Z)			
				Read(Z)		
Read(Z)						
<u>Abort</u>						
Read(W)						
		Write(W)				
<u>Abort</u>						
					Write(X)	
					Write(Z)	

Aborted process:

~~T_i R₂ (2) and w₃ (2)~~

↳ conflicting instruction older performed first

Transaction units → Transaction units
smaller timestamp larger timestamp

Validation-Based Protocol (Optimistic Concurrency Control)

Execution in 3 phases -

- 1) Read and execution phases: Transaction T_i writes only to temporary local variables.
- 2) Validation Phase - Transaction T_i performs a 'validation test' to determine if local variables can be written without violating serializability.
- 3) Write phase - If T_i is validated all updates are applied otherwise rolled back.

Validation Phase:

$TS(T_i) < TS(T_j)$ either one of the following condition holds:

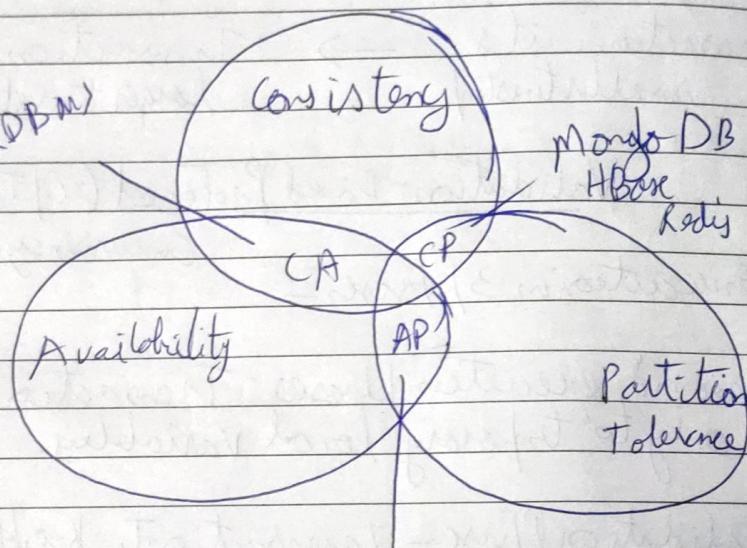
i) $finish(T_i) < start(T_j)$

ii) $finish(T_i) \leq validation(T_j)$ and set of data items written by T_i doesn't intersect with set of data items read by T_j or $validation(T_i) < finish(T_j)$

NoSQL

- Suitable for large amount of data, structured, unstructured or semi-structured.

Cap Theorem



CouchDB
Cassandra
DynamoDB

In NoSQL P has to be there, Need to select either CA.

drop A - Accept waiting until data is consistent

drop C - Accept getting inconsistent data

BASE: Basically Available, soft state, Eventual consistency.

BA: Use replication and sharding to reduce likelihood of data unavailability.

S: Soft state indicates that state of system may change over time even without input.

E: Although application must deal with instantaneous consistency, NoSQL system ensures that at some future point data assumes a consistent state.

Types:

Key Value

Column

Graph

Document

MongoDB

Open source document database designed for JSON-like data storage.

Each MongoDB instance can have zero or more databases, each acting as high-level container.

A database can have zero or more collections.

A document is made up of one or more fields.

Doc:

"address": {

 "building": "1007",

},

 "borough": "Bronx",

 "grades": [→ mult. file values

 {"date": {"\$date": "2012-04-01"}, "grade": "A",

 "score": 28}

}

creation of database → Use database name

db.createCollection("name", options)

db.createCollection("my col", {
 capped: true, →
 autoIndexId: true, size: 6142800, max: 10000})

db.restaurants.find() → Display all.

db.restaurants.find({ "borough": "Bronx" }).limit(3)

db.restaurant.find({
 "borough": "bronx",
 "cuisine": "I" }).limit(13).sort({
 "cuisine": 1 })
Display only cuisine ↗
→ ↗

Sort cuisine
based on alphabets
- 1 = DSC

Query to find restaurants with score more
than 10

db.restaurant.find({
 "grades": {
 "\$elemMatch": {
 "\$gt": 90 } } });

It takes item

Find where cuisine not american

db.restaurants.find({ "cuisine": { \$ne: "American" } })